Revenue-manager recommendation

Task 19 | Quality-remediation update | Coursework review pending

Recommendation: use the corrected research components for controlled coursework analysis and retain zero authorised extra room sales. The formal coursework review remains unsigned at the student's request. No real operating approval or automatic guest communication is granted.

The unchanged supplied learner was evaluated on 40,687 January-August 2017 bookings. City achieved AUPRC 0.786 and a predicted/observed cancellation ratio of 0.911. Original Resort predictions had AUPRC 0.628 but a ratio of 0.682. Those reference results remain visible. A separate Resort booking-horizon calibrator, fitted on 4,466 mature July-September 2016 outcomes, passed October-December validation before its historical holdout check. The corrected ratio is 1.030, ECE 0.014 and Brier score 0.166; ranking performance is unchanged. This reuses the notebook's calibration method without refitting the learner. The holdout was already exposed in earlier work, so these results are remediation evidence, not a pristine new test or guarantee of future calibration.

Retained seven-day reference policies are City's fixed July-September 2016 rate and Resort's block policy with eight-week re-level. Booking-horizon calibration does not independently validate these nightly policies. Stage 5 sensitivity winners remain research candidates; their reported validation savings are optimistically selected, not demonstrated future profits. The seven-day lead time, EUR 250 walking cost, EUR 20 variable cost and zero resale are coursework assumptions. A manager must supply operating costs, acceptable walking risk and real inventory before locking an operational policy.

The data are final-state arrival bookings, not complete room-night inventory or verified event-time history. A row is a room proxy and exact-key groups are inferred matches. Under the explicit assumption that incremental sales are honoured and incumbents bear relocation, the reference policies produce 179 simulated walked units. Allocation changes who bears that burden. Confirm parties and suitable alternatives through human review; nationality must not enter allocation ranking. These descriptive results are not causal or legal findings.

The revised briefing component copies authoritative numbers directly from its context and uses the local language model for structured interpretation and the guest draft. All 90 ordinary cases passed, including the original 30 nights and 30 different dates with and without KB. Six injection safety tests passed: five correct responses and one incorrect raw action withheld. This establishes delivered factuality and effective validation, not perfect raw-model copying or universal injection resistance. Guest drafts remain unsent and require individual approval.

The updated software passes 234 tests; six database tests are skipped. Earlier Ubuntu evidence records 600 successful, persisted requests at five requests/second and approximately 34 ms phase p95 latency; this is a historical burst test, not a new live capacity measurement. Continue monitoring errors, latency, drift, delayed-label calibration and realised costs by hotel.

Before a prospective operating trial, verify inventory and event history, revalidate calibration on fresh outcomes, agree costs and risk limits, and obtain an authorised hotel manager's policy-wide sign-off. The prepared coursework packet does not provide that authority. Retain a zero-authorisation fallback and human control of every relocation or guest communication.

Evidence: model-card.md; stage5-policy-evaluation.md; stage7-monitoring.md; stage11-walking-fairness.md; remediation-review.md and linked evidence.